



MIT SLOAN SCHOOL OF MANAGEMENT
MIT SCHWARZMAN COLLEGE OF COMPUTING



IMPLEMENTING AGENTIC AI

Your organizational playbook

Table of contents

How to use this playbook	3
Module 1: The vocabulary week	4
1.1 Define agentic AI in your own words	4
1.2 Position your organization on the continuum	5
1.3 Apply the three adoption realities	7
1.4 Identify your candidate workflow	9
1.5 Running recommendation: The WHAT	10
Module 1 rubric	12
Module 2: The building week	14
2.1 Design the agent: Five building blocks	14
2.2 Decide to build, buy, or stack	17
2.3 Set the threshold map: Human-in-the-loop	19
2.4 Draft a three-week MVP plan	21
2.5 Determine who manages this agent	23
2.6 Define AgentOps: The measurement discipline	24
2.7 Running recommendation: The HOW	26
Module 2 rubric	28
Module 3: The scaling week and board recommendation	30
3.1 Position agents on the autonomy dial	30
3.2 Determine the four safeguard layers	31
3.3 Make the comparison that matters	33
3.4 Layer lines of defense at portfolio scale	35
3.5 Manage non-human identity	36
3.6 Set differentiating governance	37
3.7 Running recommendation: The SCALE	38
Board Recommendation	40
A. Body	40
B1. Value creation across five dimensions	41
B2. ROI narrative: Near term and long term	43
C. The closing paragraph: The ask and the commitment	45
D. Executive summary	45
Module 3 rubric	48

How to use this playbook

This document is the critical learning asset of the program. You will complete it in three sittings, one per module. The choices you make in Module 1 will frame the design work you do in Module 2, which in turn will frame the scaling and governance choices you make in Module 3. By the end of the program, this document will contain a coherent, defensible recommendation for your organization, written in your voice, for your board.

Each section of this playbook corresponds to one of the three governing questions of the program:

- **Module 1:** What is agentic AI?
- **Module 2:** How will you implement it?
- **Module 3:** How will you scale it responsibly?

The “running recommendation” paragraphs in each are the through-line of the playbook. At the end of Module 3 you will assemble these, in order, as the body of your final board recommendation. Throughout, aim for the tone of an executive note: considered, candid, and written for a colleague who is smart but unfamiliar with your specific context.

Some questions provide an optional scaffold. You are not required to use this scaffold. It is offered as a starting point, in the spirit of writing a first draft you can revise. Strong paragraphs often depart from it.

Take a moment to fill in the details below. You will refer back to this Playbook throughout the program; making it yours from the outset matters.

Name	MIN QIAN
Organization	Huawei International Co. Ltd.
Role	Project Manager
Date	27 July 2026

Module 1: The vocabulary week

In this section, you make the foundational choice every agentic AI deployment depends on: deciding what you mean by the word, and where on the AI continuum your organization's first serious opportunity lives.

"Before we can tell the board what to do, we have to tell each other what we're talking about."

The fictional CEO of Boulton Industries

Module 1 provides three distinct registers: McDonagh-Smith's framing of agentic AI on the AI continuum, Jacob Andreas's technical introduction to the building blocks of an agent, and Kanchana Patlolla's practitioner perspective on what enterprise adoption looks like. Now, turn that vocabulary into a position. Not a finished plan – a position. The work is to decide, in language you can defend to a colleague, what agentic AI is, where it sits on your organization's horizon, and which workflow you intend to make the center of your learning.

1.1 Define agentic AI in your own words

"An agent is something my people are going to have to learn to manage. Not use. Manage. That's the definition I'm holding in my head."

Priya Kapoor, fictional Chief People Officer, Boulton Industries

1. In one sentence, write your own working definition of agentic AI. Aim for the sentence you would say at the start of an executive meeting in your organization.

Add your one-sentence definition here:

Agentic AI refers to an intelligent system capable of independently handling complex tasks with minimal supervision. It can autonomously decompose work steps based on high-level objectives, interact with external systems to execute tasks and achieve goals, while also possessing the ability to self-reflect and verify outcomes, continuously optimizing its behavior during operation.

2. Now reflect: What does your definition prioritize: the technical architecture, the decision authority, the unit of work, the operational queue, or the people who manage it? What does that priority reveal about how you are inclined to frame agentic AI in your organization?

Add your three to four sentences response here:

The definition prioritizes the unit of work and the technical architecture, with decision authority implicitly assigned to the system. The priority reveals that I am inclined to frame Agentic AI as an execution engine, a system that is given a task and independently figures out how to complete it. This maybe a powerful framing for engineering and product development, but it could risk overlooking the organizational and human dimensions that are often the hardest part of real-world deployment.

3. Where in your organization might colleagues currently be using the word “agent” in materially different ways from your definition? Name two specific contexts (a function, a project, a team) and describe the divergence.

Add your two specific contexts, in one short paragraph each, here:

Context 1: Colleagues in the company might refer to 'agent' as an insurance agent, rather than an AI agent. The term comes up when we need to find an agent to purchase commercial insurance.

Context 2: When we talk about chatbots, meeting assistants, and other intelligent work tools, we use the term 'agent'—typically defined as a work assistant capable of performing one specific function or a combination of multiple functions, designed to relieve human hands from handling routine tasks.

Divergence: My definition points to a complete intelligent system with goal understanding, autonomous planning, tool invocation, self-reflection, closed-loop optimization, operating with initiative, continuity, and the ability to improve over time. The "agent" in context 1 refers to a natural person or licensed entity with professional qualifications, subject to regulatory oversight, legally authorized to sign contracts, and compensated through commissions. Context 2 refers to a reactive tool, often driven by hard-coded rules or large language models, typically focused on a single function (e.g., meeting minutes generation, schedule queries). It is designed to "execute upon user command" rather than "autonomously set sub-goals and continuously optimize."

1.2 Position your organization on the continuum

“The difference between being AI-first and AI-native is the difference between a strategy and a culture, and agentic AI requires the second.”

Adapted from Kanchana Patlolla, Module 1 Unit 3

1. Locate your organization on the AI continuum. Is it primarily extracting value from classical machine learning, generative AI as a productivity tool, early agentic

deployments, or some combination of these? Be specific. Name the systems, processes, or functions that anchor your assessment.

Add three to five sentences, with concrete examples, here:

Some combination of these, but I want to share an AI-native adoption here: the company's AUTINOps Intelligent Operations solution, stands as a benchmark case of an AI-native system in the telecommunications sector. The solution condenses a previously nine-step manual operations and maintenance (O&M) workflow into a two-step intelligent closed loop. By leveraging multi-agent collaboration, a Digital Twin Network (DTN), and a domain-specific professional large model (EDNS 2.0), all built on AI-native technologies.

2. Patlolla draws a distinction between AI-first organizations (where AI appears in strategy documents and initiative lists) and AI-native organizations (where AI is the default mode of thinking and acting). Where does your organization truly sit on that spectrum? What evidence supports your assessment?

Add three to five sentences here:

The organization remains in a transitional phase, demonstrating AI-first strategy (embedding AI into existing workflows) while building toward AI-native culture (redesigning workflows around autonomous agents). However, the default mode of thinking and acting across the entire enterprise is not constituted.

Some signals indicate that native logic has not yet become the default mode of thinking and acting across the entire enterprise:

- a) *deployment is still concentrated in O&M, manufacturing, and specific telecom functions*
 - b) *many legacy processes still operate on traditional models*
 - c) *the AUTINOps transformation redefined team structures and skill requirements, but this is not yet enterprise-wide*
3. If your organization made a serious commitment to agentic AI in the next 12 months, which of the earlier waves (classical ML readiness, generative AI fluency, data infrastructure, change capacity) is the one most likely to constrain you? Why?

Add two to three sentences here:

Change capacity would be the primary constraint, because the organization's deeply ingrained hardware-centric culture and the need to reskill its massive R&D workforce pose a significant human and organizational challenge. In addition, this AI-native transformation is fundamentally constrained by the need to dismantle

data silos and unify fragmented data definitions across Huawei's product-line-centric structure.

1.3 Apply the three adoption realities

First: Speed of iteration matters more than perfection of plan. The technology landscape is moving so fast that a six-month evaluation cycle is itself a risk.

Second: The biggest barriers are organizational, not technical. Models are mostly capable enough; organizations are mostly not ready.

Third: Existing metrics will not capture the value agentic systems create. The KPIs you have today were likely built for a different unit of work.

1. **Speed of iteration:** Where in your organization is the current evaluation or procurement cadence likely to be a strategic risk for agentic AI specifically? Name a process or governance step where the cadence needs to compress, and estimate by how much.

Add two to four sentences, with a concrete cadence implied, here:

The current quarterly procurement cadence is a strategic risk for agentic AI. Agentic AI systems evolve continuously, they learn and improve weekly, not quarterly. To keep pace, the evaluation cycle needs to compress from 90 days to about 7–14 days per iteration. Treating agents like hardware procurement, rather than software that improves over time, would create a significant competitive disadvantage. A shift to continuous, outcomes-based evaluation is essential.

2. **Organizational barriers:** If your first serious agentic deployment fails or underperforms, what are the two most likely organizational reasons? Be specific: silos, incentive misalignment, data ownership, change capacity, identity issues, governance immaturity and so on.

Name two reasons, writing a short paragraph for each, here:

I think if Agentic AI deployment fails, mostly likely not because the technology is immature, but because the organization is not yet ready to let it access the data it needs or to trust it with the autonomy it requires. Both are leadership and culture problems, not engineering problems.

Reason 1: Data Silos and Fragmented Ownership

Agentic AI requires seamless, real-time access to data across previously disconnected functions, operations, supply chain, customer service, and network management and so on. However, my organization's strong product-line and business-unit structure may create intentional or unintentional barriers to cross-domain data sharing. In the organization, employees must request data access privileges based on their role and department. If the agent cannot access the full context it needs to plan and execute, its decisions will be incomplete or wrong, and the organization will blame the 'failed AI' rather than the siloed data ecosystem that constrained it.

Reason 2: Change Capacity

Agentic AI does not just automate tasks; it redefines how work is allocated between humans and machines. This threatens established roles, decision rights, and career progression paths. If middle managers or O&M engineers perceive the agent as a threat to their expertise or authority, rather than as a tool that augments their capabilities, they may passively resist, underutilize it. The over 400 engineers retrained in AUTINOps succeeded because the change was led and supported; if similar change management is not embedded from the start, the deployment will underperform regardless of technical excellence.

3. **New metrics:** Pick one workflow in your organization where, if an agent did the work, the existing metric (cycle time, cost saved, headcount avoided) would understate the value created. What new metric is missing?

Add three to five sentences, naming the workflow and the missing metric, here:

Workflow: Cross-functional incident resolution in telecom operations — where a network fault requires coordination across RAN, core network, and maintenance teams.

Existing main metric: Mean Time To Repair (MTTR), measures speed of resolution but misses the organizational cost of coordination.

Missing metric: Collaboration drag, which is the total human time spent on handoffs, escalations, and re-explaining context across teams. When an agentic AI handles resolution, it eliminates this friction by autonomously gathering data, diagnosing root causes, and executing fixes without multiple engineers interrupting their work. This drag is currently invisible in operational metrics, yet it represents a significant portion of cognitive load and delays. Tracking the reduction in cross-functional touch points would better capture agentic AI's true value.

1.4 Identify your candidate workflow

“Don’t bring the safest one. Don’t bring the most ambitious one. Bring the one that will teach us the most.”

The fictional CEO of Boulton Industries

A strong candidate workflow tends to share several characteristics. It is multi-step rather than a single prompt. It crosses functions, or at least systems. It involves judgement at one or more steps, not pure mechanical processing. It has a measurable baseline today (cycle time, error rate, cost per transaction). And, importantly, it surfaces the kind of decisions you most need to learn to make as an organization, about thresholds, safeguards, and the new role of an agent manager. The point is not to pick the workflow most likely to succeed. The point is to pick the one that will teach you the most about how to make agentic decisions well.

1. Name your candidate workflow in a single phrase, as you would on a slide. Then describe it in two or three sentences for someone outside your organization.

Add one name and two to three sentences of description here:

Candidate Workflow: AI Project Brain - Autonomous End-to-End Orchestration

AI Project Brain transforms project management from reactive administration to proactive orchestration, enabling teams to focus on value-added client engagement and strategic decision-making while the agent handles the routine tasks of contracts, budgets, and governance.

2. Where in the workflow does the human judgement currently live? Name the specific decision points and what makes each one a judgement rather than a calculation.

Add two or three decision points, in one or two sentences each, here:

Decision Point 1: Budget & Cost Monitoring

Optimization involves trade-offs between cost, speed, and quality, e.g., switching suppliers may save money but delay delivery. Human judgment is needed to assess client tolerance, relationship impact, and long-term strategy.

Decision Point 2: Client Report Generation

Reports are not just data dumps, they are relationship management tools. Human judgment determines what the client needs to hear vs. what they want to hear, and how to frame challenges constructively.

3. What is the measurable baseline today (cycle time, error rate, cost per transaction, throughput per quarter)? What numbers exist that you could compare an agent's performance against?

Add two to four numerical or quasi-numerical baselines here:

- 1) *Contract / PO Processing Cycle: Current Baseline: 3-5 business days; AI Project Brain: 0.5 business day*
- 2) *Budget Monitoring Time: Current Baseline: 4-6 hours per month; AI Project Brain: 0.5-1 hour per month*
- 3) *Meeting Minutes & Follow-Up Closure: Current Baseline: 48 hours per meeting; AI Project Brain: 2-4 hours per meeting*
- 4) *Client Report Generation: Current Baseline: 4-8 hours per report; AI Project Brain: 15-30 minutes per report*

4. What will this workflow teach your organization that an easier or safer choice would not? Be specific about the kinds of decisions it will force you to make for the first time.

Add three to five sentences here:

This workflow will force my organization to explicitly define 'when the agent's decision becomes final', setting formal thresholds for autonomous action that today exist only as human judgment calls. It will compel cross-functional data sharing agreements across Legal, Finance, and Project Ops, breaking down silos that have never been required to interoperate at the system level. It will force a new accountability framework, determining who owns errors when the agent makes a mistake, and whether the agent is a tool or a team member. Finally, it will demand a fundamental redesign of the project manager's role, shifting from administrative execution to client relationship management and strategic oversight. These are decisions that a narrow, safe automation project would never surface, precisely because they are uncomfortable, governance-heavy, and force organizational change rather than just technical deployment.

1.5 Running recommendation: The WHAT

This paragraph captures the WHAT of your recommendation. It does three things:

1. Names the strategic posture you are proposing your organization take toward agentic AI;
2. Identifies the specific workflow you intend to make the center of your initial learning;

3. Grounds both in honest reference to where your organization actually is on the AI continuum and the three adoption realities. It is the paragraph your board reads before any of the technical or governance detail.

Optional scaffold:

[Your organization] should respond to agentic AI by [a strategic posture, such as “deliberately and at measured pace”; “aggressively in two specific functions”; “through a portfolio of focused experiments” etc.], focusing initial learning on [your candidate workflow], because [your honest assessment of where the organization is, and what the three adoption realities imply for it].

1. Write a coherent WHAT paragraph. Do not use bullet point. This text will sit at the top of a board document.

Add four to seven sentences here:

Huawei should respond to agentic AI by deliberately and at a measured pace, focusing initial learning on the AI Project Brain, autonomous end-to-end orchestration from contract to delivery.

This measured approach fits Huawei's current position: strong in classical ML and generative AI, but still transitioning from AI-first to AI-native, with promising pockets of native deployment (e.g., AUTINOps), yet lacking the unified data infrastructure and cross-functional governance needed for safe, large-scale agentic autonomy. The three adoption realities make this clear: agentic AI requires breaking data silos, establishing new accountability for machine errors, and fundamentally redefining human roles. By starting with a contained, high-visibility workflow that touches commercial, financial, and delivery functions without yet impacting customer-facing or network-critical systems, Huawei can safely build the organizational muscle for data sharing, governance, and role redesign, generating the evidence and trust needed to scale agentic AI confidently, without overcommitting before the cultural and infrastructural foundations are ready.

2. Read your paragraph back. If a member of your board read only this paragraph and nothing else, what would they think you are recommending? Is that what you want them to conclude? Revise accordingly.

Add a short revision note to yourself, if any, here:

Board need to hear both: ‘we are deliberate’ and ‘we are ambitious’. To make clear that this is not a small experiment, instead it is the first, high-signal move in a sequenced transformation, designed to unlock the entire organization's agentic future.

'Huawei should respond to agentic AI with deliberate, strategic urgency, anchoring its initial bet on the AI Project Brain, an end-to-end autonomous workflow that spans contract management, budget control, and governance reporting.'

Module 1 rubric

	Criteria not met	Criteria met	Good	Exceptional
Sections 1.1–1.3 <i>Vocabulary and context</i>	No submission. OR The response fails to define agentic AI or to place it on the AI continuum.	The response attempts to define agentic AI and to place it on the AI continuum, but insight into the three adoption realities are vague.	The response provides sound insight, defining agentic AI, placing it on the AI continuum, and applying the three adoption realities.	The response provides thorough insight, defining agentic AI clearly, placing it on the AI continuum, and rigorously applying the three adoption realities. The response demonstrates a high level of expertise in the module's content.
Section 1.4 <i>Candidate workflow</i>	No submission. OR The response does not demonstrate insight into the candidate workflow.	The response attempts to provide insight into the candidate workflow.	The response provides sound insight into the candidate workflow.	The response provides thorough insight into the candidate workflow and demonstrates a high level of expertise in the module's content.
Section 1.5 <i>Running recommendation: WHAT</i>	No submission. OR The response does not provide a	The response attempts to provide a sound running recommendation.	The response provides a sound running recommendation.	The response provides an insightful running recommendation and demonstrates

	Criteria not met	Criteria met	Good	Exceptional
	sound running recommendation.			a high level of expertise in the module's content.
<i>Clarity and logic of the submission's structure</i>	No submission. OR The response is unstructured, unclear, and lacks succinctness.	The response is logically structured, clear enough to comprehend, and mostly succinct.	The response is well structured in terms of logic, clarity, and succinctness.	The response is exceptionally well structured in terms of logic, clarity, and succinctness.

Module 2: The building week

In this section, you turn the candidate workflow you identified into a designed agent, making the architectural, adoption, and operational decisions that determine whether what you build will work in the world.

“An agent is a language model plus scaffolding.”

Adapted from Jacob Andreas, Module 2 Unit 1

2.1 Design the agent: Five building blocks

Andreas argues that building a language model agent requires answering five questions, and that every design decision downstream is a commitment on one of these five axes. The discipline of working through them in sequence is what separates a real agent from a wished-for one.

“What turns a language model into an agent is not a new kind of model but a change in interpretation: The same text output, interpreted as a program rather than as a reply, becomes an action.”

Adapted from Jacob Andreas, Module 2 Unit 1

Complete each row of Table 1 for your candidate workflow. Be concrete. Where you do not know the answer yet, name the open question – that is itself a useful output.

Table 1: Five building blocks.

Building block	The question to answer	Your design decision
Perception	How will the agent observe the world? What inputs, in what formats, from what sources?	The agent observes the world through Organization’s internal operational systems: MetaERP for finance and budgets; ISDP for project tasks, milestones, and resource scheduling; iSales for contract repository; email and Teams Meeting for communication; and external benchmarks for market context. Inputs arrive as structured API data, unstructured text from emails and meeting transcripts, and dependency

		graphs etc.
Action	What is the agent allowed to do? Which tools, APIs, or systems can it invoke, and with what side effects?	The agent is allowed to execute a set of actions based on predefined thresholds. It can autonomously update task statuses, adjust resource allocations, and generate progress reports by calling ISDP's API. It can query financial data, flag budget variances, and release payments only below a configurable monetary limit (e.g., \$5,000) by invoking MetaERP; any larger financial action requires human approval. In iSales, the agent can draft change orders and generate contract summaries, but cannot execute legally binding signatures. It can draft email and schedule meetings by calling email and Teams Meeting, but all external client-facing communications require human review before sending. All actions are fully logged for audit and rollback capabilities.
Planning	How will the agent reason about multi-step tasks? Single-step, chain-of-thought, decomposition into sub-agents?	For multi-step tasks, the agent employs a hybrid reasoning strategy tailored to complexity. For routine, well-defined actions like updating a task status in ISDP, it may use single-step execution. For complex decisions such as resolving a budget overrun, it must apply chain-of-thought reasoning, explicitly breaking the goal into a logical sequence: analyze variance, identify root causes, simulate corrective scenarios, and propose actionable recommendations. For highly intricate, cross-functional workflows that span multiple systems, for instance, aligning ISDP task adjustments with MetaERP financial approvals and iSales contract

		impacts, the master agent decomposes the work into specialized sub-agents, each handling a specific domain, meanwhile the master agent is responsible for managing handoffs, conflicts, and ensuring end-to-end coherence. Before executing any critical action, Human-in-the-loop is required, the agent always surfaces a human-readable rationale for review and approval.
Memory	What does the agent need to remember between turns, between sessions, between deployments?	The agent maintains multi-tier memory: firstly, within-session (current project context, pending decisions), secondly, between-sessions (historical project data, team preferences, past resolutions), thirdly, between-deployments (institutional knowledge, approved heuristics, vetted decision patterns stored in a version-controlled knowledge base). Memory is actively curated: sensitive data perhaps is auto-purged per retention policies, and each deployment triggers a deliberate review to retire obsolete patterns, ensuring the agent learns without past mistakes or biases.
Safety	How is risk architected into the surrounding system, not the model? What does the system do that the model cannot be trusted to do?	Risk is architected into the surrounding system through hard guardrails that the model cannot enforce. The system, not the model, enforces hard financial caps (e.g., blocks MetaERP payments > \$1,000), mandatory escalation routing (e.g., forces PMO review for strategic accounts), and output validation (e.g., rejects malformed ISDP task updates). It maintains immutable audit logs and enables rollback of any action. The model proposes, but the system disposes verifying, constraining, logging, and reversing

		actions.
--	--	----------

1. Andreas posits that you should not rely on the language model itself to manage risk. Instead, you architect the safety into the surrounding system. Looking at your Safety row above, where in your design is the riskiest place you are currently relying on the model itself to make a judgement that it shouldn't?

Add one specific place and one or two sentences on why moving it outside the model would matter:

The riskiest place relying on the model itself could be in the Planning: Decomposition into Sub-agents step, specifically, when the master Agent interprets a vague project goal like "fix the budget overrun" and autonomously decides which sub-agents to call and in what order, without external validation.

Why moving this outside the model matters: The model may hallucinate a non-existent dependency (e.g., call iSales before MetaERP) or omit a critical compliance step, leading to an irreversible workflow chain. By moving the decomposition logic into a version-controlled, human-approved workflow manifest that the system interprets rather than the model generating it on the fly to ensure that the overall orchestration follows auditable, tested paths, while the model only fills parameters and handles exceptions within that bounded framework.

2.2 Decide to build, buy, or stack

Patlolla addresses the strategic posture question every leader will face: At each layer of the agentic stack (model, framework, platform), do you build, buy, or stack a combination? Her answer is consequential and disciplined. The decision is driven by the autonomy you need and the skill set you have, not by technology preference.

"It doesn't have to be one or the other. It could be a combination of all three in your enterprise."

Kanchana Patlolla, Module 2 Unit 2

For your candidate workflow, complete the build/buy/stack matrix in Table 2. In each cell, note the choice you are inclined toward, and one or two sentences justifying it for this specific deployment.

Table 2: Build, buy, stack matrix.

Layer	Build	Buy	Stack (combination)
Model		Data cannot leave the internal network, so public APIs are ruled out. We will procure a commercial model that supports private deployment running on Huawei's own infrastructure. Building a foundation model from scratch is not viable.	
Framework			Use an open-source orchestration framework (e.g., LangGraph, AutoGen) in an offline, internal version as the base. Build a proprietary governance and safety wrapper on top to enforce internal approval matrices, compliance policies, and audit trails, this is where organizational control is critical and cannot be outsourced.

Platform	Leverage Huawei's internal AI platform for compute and model serving. Build a custom business abstraction layer to unify and simplify access to ISDP, MetaERP, iSales, especially to compensate for ISDP's poor user experience, with all data flows strictly confined to the internal network.		
-----------------	---	--	--

1. Patlolla's underlying argument is that the choice at each layer is a function of two variables: the autonomy you need, and your internal skill set. For your candidate workflow, which layer is most likely to expose a skill-set gap in your organization? What would closing that gap look like (building internal capability, hiring, partnering, etc.)?

Add three to four sentences here:

The Framework layer is most likely to expose a skill-set gap, particularly the potential self-built governance and safety wrapper components. This requires deep expertise in policy engineering, compliance rule encoding, and audit trail design, which differs from traditional software development or IT operations. Closing this gap would require internal capability building (e.g., upskilling existing security and compliance teams in rule-as-code and agent safety frameworks) combined with targeted hiring of specialists in AI governance and safe agentic systems. The long-term ownership must reside in-house due to Organization's stringent data security and regulatory requirements.

2.3 Set the threshold map: Human-in-the-loop

"I don't care about a \$5 fraud, let it go. That's something we've always done, even pre-generative AI: We've created thresholds and handled those cases."

Kanchana Patlolla, Module 2 Unit 2

For your candidate workflow, draw the threshold map. In Table 3, identify three decision types the agent will encounter. For each, name the threshold at which the agent acts autonomously, the threshold at which it escalates, and the threshold above which a human must originate the decision.

Table 3: Threshold map.

Decision type	Agent acts autonomously when...	Agent escalates to human when...	Human originates the decision when...
1. Project Task & Resource Adjustment	Updates non-critical subtasks, logs hours, reallocates low-priority resources within a single project without changing milestones.	Changes involve critical milestones, cross-project resource conflicts, or external supplier requests.	Pausing/terminating a project, changing core scope or budget baseline.
2. Budget & Financial Approval	Queries budget usage, generates forecasts, flags variances	Releases funds over \$1,000 or adjustments involving loss-making projects or high-risk client tags.	Any release above \$5,000, creating new budget lines, or modifying the overall project budget baseline.
3. Contract & Sales Change	Retrieves contract clauses, generates standard NDA drafts, flags upcoming renewals.	Drafts change orders, modifies non-core terms (e.g., payment cycles), or handles non-standard requests from strategic clients.	Signing any legally binding documents, modifying core commercial terms (price, deliverables, liabilities), or terminating contracts.

1. Looking at your threshold map, which threshold are you least confident about (where you genuinely do not yet know where the line should be)? What experiment would help you find out?

Add one specific threshold and one or two sentences on the experiment that would calibrate it here:

The threshold I am least confident about is the contract change order escalation range. While we do have historical contract change records, they lack decision context (e.g., change rationale, client relationship weight) and contain no 'rejected' negative samples, making it difficult to fully understand real risk exposure.

Calibration method: First, clean and analyze the last 2–3 years of historical data to build a baseline statistical distribution (e.g., correlation between change amount, approval tier, and downstream disputes) to set an initial threshold. Then, during a pilot, have the agent log (but not act on) all contract change requests, tracking human decisions and their outcomes. After three months, compare actual observed risk against the historical baseline and iteratively adjust the threshold.

2.4 Draft a three-week MVP plan

Patlolla's recommended cadence is three weeks (and five people) from design board to first internal release. The discipline is not in the speed but the scope. A three-week MVP is honest only when the scope is held tight. Pick one slice of your candidate workflow that can credibly be agentified in three weeks.

"Three weeks is even a lot, because you can build a prototype for a complex vision in less than a weekend."

Kanchana Patlolla, Module 2 Unit 2

1. **MVP scope:** What is the smallest, end-to-end slice of your candidate workflow that an agent could plausibly handle? Describe it the way you would describe it to the engineer who would build it. Be specific about start and end points.

Add two to four sentences here:

We are building a read-only, single-project agent that runs daily at 2:00 AM. It queries ISDP for tasks with actual finish dates deviating ≥ 5 working days from baseline, reads the last 7 days of Teams/email/WeLink logs for that project to extract two root cause labels, and generates a structured deviation report with exactly two corrective recommendations. The agent does not execute any action,

it only sends the report via WeLink/email to the assigned PM and creates a follow-up task requesting the PM to accept, reject, or ignore the suggestions.

2. **MVP audience:** Who are the five (or fewer) people who would use the MVP first? Name the role and describe the type of feedback you would solicit from each.

Add up to five named roles, one line each, here:

1. Project Manager (End User): Does the deviation report save you time? Are the root-cause labels and recommendations accurate and actionable?
 2. ISDP Product Owner (Data & System Owner): Are the APIs providing the data we need in the right format and frequency? Are there data quality issues affecting the agent's performance?
 3. IT Security & Compliance Lead (Risk Approver): Does the agent's read-only access and internal data flow meet our security policies? What additional safeguards or audit logs are required before we scale
 4. AI Engineer (Technical Implementation): Is the current prompt and reasoning logic performing as expected? Can we trace each decision to the right data source?
3. **What to leave out:** The hardest part of a three-week MVP is what you leave out. Name two specific aspects of your candidate workflow that you are deliberately deferring beyond the MVP, and the trigger condition that would bring each one back into scope.

Add two deferred aspects and triggers, in one or two sentences for each, here:

Deferred Aspect: 1. Cross-project resource conflict detection

Trigger Condition: Bring this back when the single-project deviation MVP is validated as useful, and PMs explicitly request visibility into whether adjusting one project's resources will impact others.

Deferred Aspect: 2. Automated corrective actions (e.g., automatically reassigning tasks in ISDP or releasing budget in MetaERP)

Trigger Condition: Bring this back when the agent's recommendation accuracy (PM adoption rate) remains stable at $\geq 85\%$ for four consecutive weeks, and IT Security has completed a risk assessment for write-level system access.

2.5 Determine who manages this agent

“Every agent needs a manager.”

Adapted from Kanchana Patlolla, Module 2 Unit 2

1. Who, by name or by role, will be the day-to-day manager of your agent? Why this person? What do they bring that the engineering team would not?

Add three to five sentences here:

The day-to-day manager of the agent will be a Project Manager (PM) from the pilot team. The PM is the direct end-user and the most qualified judge of whether the agent's reports are actionable and accurate; they know what constitutes a useful recommendation in real project contexts. Unlike the engineering team, the PM can provide qualitative feedback on why a suggestion was ignored or rejected, which is critical for iterating on the agent's reasoning logic. Engineering can optimize performance and stability, but only the PM can validate whether the agent is actually delivering business value, making them the essential ‘human in the loop’ for this workflow.

2. What does the agent-manager role need to do that this person's current role does not? Name three specific responsibilities that are new (not adapted versions of existing ones).

Add three specific new responsibilities, in one short paragraph each, here:

1. Review and Accept/Reject Agent Recommendations

The PM currently performs deviation analysis and makes decisions themselves. The new responsibility is to review the agent's daily deviation report and explicitly accept, reject, or ignore each recommendation, including a brief reason for rejection (e.g., "root cause misidentified," "recommendation not feasible"). This feedback becomes the critical training data for the agent's continuous improvement, something the PM never had to provide before.

2. Maintain and Refine Threshold Rules

The PM currently does not define system rules. The new responsibility is to participate in monthly reviews of the agent's autonomy and escalation thresholds, for example, is "trigger report at ≥ 5 working days delay" too sensitive? Based on real business experience, the PM recommends threshold adjustments to the PMO/IT team. This elevates the PM from a rule executor to a rule co-designer.

3. Identify and Flag Edge Cases

The PM currently handles only their own projects and does not proactively flag a typical cases. The new responsibility is to actively flag cases where the agent fails or provides incorrect recommendations as "edge cases," with a brief explanation. These edge cases become essential inputs for the next phase of agent capability expansion and help engineering teams refine reasoning logic, a form of institutional knowledge capture the PM was never asked to do before.

3. Will this be an add-on to an existing role, or a new role? If an add-on: How will their existing workload be adjusted? If new: What is the path to formalizing this position in your organization?

Add two to four sentences here:

During the pilot, the agent-manager role remains an add-on to the existing PM role, as uncertainty is high and organizational change should be minimized. After the pilot, the decision depends on the real situation: agent adoption rate, actual time saved per PM, and the organization's scale-up plan. After tradeoffs, we can formalize the role either by updating the Senior PM job description or by creating a dedicated "AI Project Operations Specialist" position, but this role is not an external hire, it's from an internal promotion path for experienced PMs, turning deep domain expertise into AI governance capability, and incentivizing PMs to grow into a new, higher-value role.

2.6 Define AgentOps: The measurement discipline

Patlolla's second high-conviction framing is the move from CI/CD/CT/CM (continuous integration, delivery, training, monitoring – the MLOps inheritance) to what the field is calling AgentOps. The shift is to continuous evaluation as a first-class deployment concern, with skew detection built into the architecture.

"The dashboard is not a report. It is a conversation we are now having continuously with our own system."

Rebecca Holloway, fictional CFO, Boulton Industries

1. **Continuous evaluation:** What does the agent need to be evaluated against, in production, on a continuous basis? Name two metrics that should be tracked in real time, and one quality measure that requires periodic human review.

Add two real-time metrics and one quality measure with a proposed review cadence here:

Real-Time Metrics (Automated, Dashboard-Viewable):

- 1) Recommendation Adoption Rate: Percentage of agent-issued recommendations that are explicitly accepted (vs. rejected or ignored) by the PM within a defined time window
- 2) Escalation-to-Action Latency: Time elapsed between the agent's "escalate to human" trigger and the PM's first action on that escalated case.

One Quality Measure: Root Cause Label Accuracy

Bi-weekly: The Agent Manager reviews a random sample of 20 deviation reports, comparing the agent's root-cause labels against their own judgment, and flags mismatches. This identifies whether the agent's classification logic is drifting and informs whether the label taxonomy or prompt needs updating.

2. **Skew detection:** Where in your candidate workflow is the agent's behavior most likely to drift over time? Is it because the input distribution shifts, the underlying model is upgraded, or a tool's API changes? What would early warning of that drift look like?

Add two to four sentences here:

Most Likely Drift Point: Root-Cause Label Classification

The root cause of drift is highly likely input distribution shifts, not the underlying model or tool APIs. As new projects roll in and PM teams evolve, the language used in Teams Meeting and email logs changes naturally. For example, a PM might start using "supplier bottleneck" instead of the previously common term "supplier delay". Because the agent learned its label taxonomy from historical data, such shifts in terminology can degrade its classification accuracy over time.

Early warning is a sustained ≥ 10 -point drop in recommendation adoption rate over a 2-week rolling window, signaling that the agent's labels are no longer matching current project language and that a label review or prompt update is needed.

3. **Ownership:** Who reviews the AgentOps dashboard, on what cadence, and with what authority? This is a governance question as much as an engineering one.

Add three to five sentences here:

The AgentOps dashboard is reviewed daily by the agent manager to detect anomalies and flag emerging issues; the PMO Lead reviews it bi-weekly to approve threshold adjustments and process improvements; and IT Security

conducts a monthly audit to ensure compliance and holds the authority to freeze the agent immediately if a violation is detected.

All anomaly resolutions are documented and fed back into the knowledge base to prevent recurrence. This ensures that the dashboard is not a human-monitoring tool but an early-warning system that triggers clear, pre-defined escalation paths and ownership.

2.7 Running recommendation: The HOW

The HOW paragraph integrates the work of this section into a single coherent statement. It tells the reader what kind of agent you are building (the building-block specifications, in compressed form), what your build/buy/stack posture is, where the human-in-the-loop thresholds sit, who will manage the agent day-to-day, and what AgentOps discipline will hold it to standard. Five things in one paragraph.

Optional scaffold:

[Your organization] will build this capability by [your build/buy/stack posture], deploying an agent with [a one-line summary of the building-block decisions] and human-in-the-loop thresholds at [the consequential decision points], managed day-to-day by [the agent-manager role], and held to standard through [your AgentOps cadence].

1. Write your HOW paragraph (again, do not use bullet points).

Add five to eight sentences here:

*Huawei will build this capability by buying and deploying a private, on-premises LLM, stacking open-source orchestration frameworks with a self-built governance and safety wrapper, and layering a custom business abstraction layer on top of its internal systems (ISDP, MetaERP, iSales, WeLink, etc.) to shield the agent from legacy complexity. The agent will use hybrid reasoning, single-step for routine actions, chain-of-thought for complex deviations, and sub-agent decomposition for cross-functional workflows with **perception** spanning structured APIs and unstructured meeting/email logs, **memory** maintained across sessions and deployments, and **safety** enforced via hard financial caps, mandatory escalation routing, and immutable audit trails. Human-in-the-loop thresholds are defined by a threshold map where the agent acts autonomously for low-risk tasks, escalates to PMs for medium-risk decisions, and requires human origination for high-risk irreversible actions. The agent will be managed day-to-day by a Project Manager during the pilot, transitioning post-pilot to an AI Project Operations Specialist, an internal promotion path from experienced PMs, who will maintain thresholds, curate edge cases, and train new PMs. It will be held to standard through an AgentOps discipline with real-time metrics*

(recommendation adoption rate, escalation-to-action latency), monthly human audits of root-cause label accuracy, and a governance cadence where the Specialist reviews daily, the PMO Lead approves bi-weekly, and IT Security holds veto power for compliance violations.

2. Read your two running recommendation paragraphs together. Do they cohere? If a colleague read both as a single argument, would they understand how the design choices in the second paragraph follow from the strategic posture in the first? Note any revision your earlier paragraph now needs in light of the work you have just done.

Add a short revision note to yourself, if any, here:

The two paragraphs seems coherent. The design decisions in Paragraph 2 are direct responses to the strategic constraints identified in Paragraph 1 (system heterogeneity, data silos, governance gaps, the need for incremental organizational learning).

I think the only revision needed is to add ISDP's user experience limitation to Paragraph 1 as an explicit justification for the business abstraction layer in Paragraph 2. The first paragraph does not currently mention the ISDP user experience challenge. Since ISDP is "not very user-friendly," this is a key driver for the "self-built business abstraction layer" in Paragraph 2. This should be reflected in the first paragraph's description of data/system complexity or integration friction.

Module 2 rubric

	Criteria not met	Criteria met	Good	Exceptional
Sections 2.1, 2.2, & 2.4 <i>Design and architecture</i>	No submission. OR The response fails to complete the five building blocks table, the build, buy, or stack matrix, or the three-week MVP plan.	The response attempts to complete the five building blocks, the build, buy, or stack matrix, and the three-week MVP plan.	The response provides sound insight on the five building blocks, a clear build, buy, or stack approach, and a realistic three-week MVP plan.	The response provides thorough insight on the five building blocks, a clear build, buy, or stack strategy, and a specific three-week MVP plan. The response demonstrates a high level of expertise in the module's content.
Sections 2.3, 2.5, & 2.6 <i>Operations and management</i>	No submission. OR The response does not demonstrate insight into the thresholds and management plans.	The response attempts to provide insight into the thresholds and management plans.	The response provides sound insight into the thresholds and management plans.	The response provides thorough insight into the thresholds and management plans and demonstrates a high level of expertise in the module's content.
Section 2.7 <i>Running recommendation: HOW</i>	No submission. OR The response does not provide a sound running recommendation.	The response attempts to provide a sound running recommendation.	The response provides a sound running recommendation.	The response provides an insightful running recommendation and demonstrates a high level of expertise in the module's content.

	Criteria not met	Criteria met	Good	Exceptional
<i>Clarity and logic of the submission's structure</i>	No submission. OR The response is unstructured, unclear, and lacks succinctness.	The response is logically structured, clear enough to comprehend, and mostly succinct.	The response is well structured in terms of logic, clarity, and succinctness.	The response is exceptionally well structured in terms of logic, clarity, and succinctness.

Module 3: The scaling week and board recommendation

“Scaling is not enlarging a pilot. It is transitioning from a controlled environment to an operational context that no one has fully tested.”

Paul McDonagh-Smith, Module 3 Unit 1

Module 3 shifts the level of analysis from the agent to the system. In this section, you complete the work of the program: positioning your agent on the autonomy dial, designing the safeguard architecture that makes it durable, and articulating the governance posture that will let your organization deploy not one agent but many – responsibly, measurably, and over time.

3.1 Position agents on the autonomy dial

“Autonomy is not a switch you flip. It is a dial you calibrate.”

Paul McDonagh-Smith, Module 3 Unit 1 Video 2

The critical design decision is where you position each agent, for each task, in each context:

- **Constrained:** Agent only returns pre-verified responses.
 - **Bounded:** Agent operates within a narrow, validated decision space with near-deterministic rules.
 - **Supervised:** Agent acts but every consequential action is reviewed before execution.
 - **Delegated:** Agent acts autonomously within thresholds; humans review by exception.
 - **Autonomous:** Agent acts and self-monitors; humans review only on flagged anomalies.
1. For your candidate workflow, determine autonomy for each of the three core decision types in your Module 2 threshold map and explain your placement in one or two sentences each.

Add three placements with reasoning here:

1. Project Task & Resource Adjustment (ISDP) Level: Supervised

Task updates and resource changes can impact milestones and dependencies. The agent proposes adjustments, but a PM must approve any action affecting critical paths or cross-project resources to maintain human control over delivery outcomes.

2. Budget & Financial Approval (MetaERP) Level: Delegated

Low-risk releases (e.g., below \$1,000 with cumulative limits) run autonomously. Any release above \$5,000 or involving high-risk accounts requires human origination.

3. Contract & Sales Change (iSales / 360 Contract) Level: Constrained / Supervised

Standard NDA generation and clause retrieval are constrained. Any contract modification or change order must be supervised: the agent drafts and flags risks, but legal/commercial sign-off remains human-only to enforce hardcoded constraint layer.

2. Where on the dial does the agent currently sit by default in your design? Where do you imagine it sitting in 12 months, if it performs as intended? What conditions would have to be met for it to migrate from the first position to the second?

Add three to five sentences here:

The agent will start at Delegated, autonomous within strict thresholds (\$1,000 budget caps, non-critical ISDP updates), with humans reviewing by exception via the AgentOps dashboard. In 12 months, if performance holds, it will evolve to a dynamic Delegated/Supervised mix, where routine tasks remain autonomous but higher-stakes decisions (e.g., resource conflicts, change orders) require conditional human approval.

I think migration shall require four conditions: $\geq 90\%$ recommendation accuracy for four weeks, escalation rate below 10%, $\geq 80\%$ PM confidence, and proven audit/rollback effectiveness. The success metrics are merely preliminary ideas. The dial moves not by time, but by evidence.

3.2 Determine the four safeguard layers

For your candidate workflow, complete the four-layer safeguard in Table 4. Each row asks you to translate the architectural principle into a safeguard specific to your design. Where a layer does not yet exist in your design, name what it would take to add it.

Table 4: Four-layer safeguard.

Safeguard layer	Architectural principle	Concrete safeguard in your design
Hard-coded constraints	Non-negotiable boundaries built into the system architecture itself, executing independently of the model's reasoning. The model cannot talk its way past them.	Financial hard cap: System-level rule blocks any MetaERP payment over \$5,000 regardless of agent output. Read-only credentials: Default agent API key has no write access to contract signature fields or budget baselines, separate human-authenticated credentials required for changes.
Output verification	Before an agent's output reaches a user or triggers a downstream action, it is checked against formal rules, consistency criteria, or a second model acting as reviewer. The two-person rule, in software.	Consistency check: For any ISDP task update, the system verifies that the task ID exists and the new dates are within permissible boundaries before passing the action downstream.
Reversibility by design	Wherever possible, agent actions are designed to be reversible. Reversibility limits the blast radius of an error; it does not prevent the error itself.	Two-stage budget release: Any agent-initiated budget release (up to \$5,000) enters a "pending" state for 2 hours before execution; PMs can cancel within that window via the AgentOps dashboard. Soft task updates: Changes to ISDP task assignments or dates are initially applied to a "draft plan" and only synchronized to the production schedule after PM confirmation or a 24-hour timeout (if no objection).
Escalation architecture	Every agent has a clear, tested escalation path. This is a mechanism for recognizing the boundary of its confidence or its authority and routing the decision to a human.	Threshold-based: Payments >\$5,000, contract changes, and strategic decisions auto-route to designated human with full context package. Confidence-based: Recommendations below 70% confidence flagged for review

		with top uncertainty factors listed. 48-hour overflow: No human response within 48 hours triggers escalation to PMO Lead with backup reviewer added.
--	--	--

1. Look across your four layers. Which is the strongest in your current design? Which is the weakest? What single intervention would most improve the weakest layer?

Add three or four sentences here:

Hard-coded constraints and escalation architecture are already robust; the weakest part is reversibility, because contract drafts and cross-system updates are difficult to undo once executed. The weakest layer is reversibility by design, as cross-system rollbacks (ISDP, MetaERP, iSales) are technically complex and often impractical once actions are externalised.

The single most effective intervention is "pre-execution staging + atomic commit" : all agent actions write to an isolated staging area; a human gate halts any external facing action (email, payment, contract draft); approval atomically commits to production, while rejection or timeout simply discards the staging data.

3.3 Make the comparison that matters

When evaluating the risk of an agentic system, the comparison is not agent versus perfection, it is agent-with-safeguards versus the current process. The question to put to your board is not "Does the agent make errors?" but "Does this agent, with these safeguards, produce better outcomes than the process it replaces, on the dimensions that matter?"

"We are replacing processes which normally humans do, and they also make mistakes."

Tim Kraska, Module 3 Unit 2

1. Sketch the head-to-head comparison for your candidate workflow. On the one hand, the current human-driven process (accuracy, cycle time, consistency, audit trail, error patterns, cost). On the other, the agent-with-safeguards. Where the projection is uncertain, say so.

Add four to six sentences, or a structured comparison, here:

Head-to-Head Comparison: AI Project Brain vs. Current Human Process

1. Accuracy

Current human-driven process: PMs manually scan ISDP for deviations; misses occur under high load or when managing multiple projects.

AI agent-with-safeguards: Agent scans all active projects daily at 2:00 AM, detecting every deviation ≥ 5 working days. Root cause labels may drift with terminology changes (e.g., "supplier bottleneck" vs. "supplier delay").

2. Cycle Time

Current human-driven process: PMs compile status reports weekly or bi-weekly; 3-5 days from deviation occurrence to notification.

AI agent-with-safeguards: Agent generates and sends reports within 24 hours of detection (daily automated scan).

3. Consistency

Current human-driven process: Thresholds vary by PM (some flag 5-day delays, others 10-day); report formats and quality differ.

AI agent-with-safeguards: Unified threshold (≥ 5 days), standardized report format, identical reasoning every time.

4. Audit Trail

Current human-driven process: Records scattered across emails, meeting notes, and Excel files—difficult to trace decisions.

AI agent-with-safeguards: Every report, root cause label, data source, and PM feedback is logged immutably.

2. On which dimension is the agent likely to be worse than the current process, at least at first? How will you defend that trade-off to your board, and what is the path to closing the gap?

Add two to three sentences here:

The agent will likely be worse on root-cause classification accuracy, as PMs bring tacit knowledge (client history, team dynamics) that the model cannot infer from meeting logs. I would defend this trade-off by framing it as a small, visible, and correctable error rate against major gains in speed, consistency, and auditability that the current process lacks. The gap will close through a continuous feedback

loop: PM corrections feed into periodic prompt/retraining updates, with improvement tracked on the AgentOps dashboard.

3.4 Layer lines of defense at portfolio scale

Most organizations that scale agentic AI successfully will not deploy one agent; they will deploy ten, or twenty, or a hundred, each calibrated differently. At that scale, governance cannot be designed agent-by-agent. It has to be infrastructural.

1. **Foundation layer: Data and model governance.** What policies, processes, or infrastructure does your organization already have in place that would extend to a portfolio of agents, and where are the most significant gaps? Be specific.

Add three to four sentences here:

The existing ISO/IEC 42001-certified AI management system and mature enterprise data governance framework provide a solid foundation that can extend to a portfolio of agents. The most significant gap is Non-human Identity Governance. The current IAM systems lack full lifecycle management for agents (registration, dynamic least-privilege permissions, and cross-system audit trails). This gap should be closed by designing and implementing agent identity lifecycle processes, aligned with my organization's existing data security and AI safety frameworks. This work is a prerequisite before scaling beyond the initial MVP.

2. **Operational layer: Runtime guardrails.** What input/output moderation, prompt routing, and use-case-specific constraints would need to be standardized across your agent portfolio? Who would own them?

Add four to six sentences here:

For the AI Project Brain portfolio, runtime guardrails must be standardized at the agent orchestration layer, not re-implemented per use case. This includes input/output moderation, such as screening for PII, toxic language, or policy violations before they reach the model or external systems; prompt routing, such as directing each request to the appropriate agent/sub-agent based on intent and context; and use-case-specific constraints (e.g., ISDP task update boundaries, MetaERP financial caps, iSales contract clause restrictions). These guardrails should be owned by a central AgentOps team, including the AI Project Operations Specialist, an IT security representative, and a PMO governance lead. Individual business units (PMO, Finance, Legal) retain the right to propose domain-specific thresholds, but the central team approves and enforces them to ensure consistency, auditability, and compliance across all agent deployments. This prevents fragmentation and ensures that each new agent inherits a proven safety baseline.

3. **Transparency layer: The AI bill of materials.** For your candidate workflow, sketch the components that would appear on its AI bill of materials (training data provenance, model architecture, version history, third-party components, tool dependencies). Where is the documentation thin?

Add five to seven sentences, or a structured list, here:

The AI Bill of Materials for the AI Project Brain would include the private LLM (model name, version, hash), fine-tuning artifacts/adapters, system and task-specific prompts (version controlled), tool/API dependencies (ISDP, MetaERP, iSales, 360 Contract, WeLink with endpoint versions), agent identity and permission records, and the underlying software stack (LangGraph, Python libraries).

The thinnest area should be the data provenance and permissions documentation. We could lack a systematic record of which historical ISDP/Teems datasets were used for prompt calibration, their versioning, and their business sensitivity classification. Similarly, the agent's cross-system permissions aren't currently documented as part of a formal asset inventory.

3.5 Manage non-human identity

1. For your candidate workflow, define the agent's identity. What is the agent allowed to access? What is it explicitly not allowed to access? Which tools, APIs, and systems does it have permissions for? What is the principle of least privilege, applied here?

Add five to seven sentences here:

The agent, named AI Project Brain Agent, is a non-human service account with a unique ID and independent audit trail, distinct from any human user. It is permitted to read task/resource data from ISDP, budget usage from MetaERP, contract metadata from iSales/360 Contract, and meeting logs from Teems/email via approved APIs. It is explicitly denied access to HR systems (employee PII/salary), customer-sensitive commercial data (unless de-identified), and any write operations on budget baselines, contract signatures, or production task schedules outside the human-gated staging workflow.

The principle of least privilege is enforced through read-only default access; every write or external action requires a separate human-approved execution channel, with cross-system operations coordinated via the orchestration layer. This ensures the agent touches only the minimum data required for deviation detection, while all high-impact actions are routed through the escalation and safeguard layers.

2. How will the agent's permissions be reviewed, by whom, and on what cadence? Under what conditions would you revoke or restrict permissions? Treat this as you would when considering a new senior employee.

Add three to five sentences here:

Permissions will be reviewed quarterly by the AI Project Operations Specialist, with IT Security and PMO support, and additionally whenever the agent's adoption rate drops by $\geq 10\%$ over two weeks or a security flag is raised. Revocation or restriction triggers include: three consecutive failed monthly accuracy audits, policy/regulation changes, sustained anomalous behavior (e.g., repeated blocked write attempts), or a manager's explicit request. Restrictions (e.g., read-only downgrade, API blocking) are documented as formal asset records in the AI BOM, with restoration requiring the Specialist's review and PMO approval. This mirrors how we would manage the permissions of a senior employee: regular reviews, clear triggers for removal, and a documented approval path for reinstatement.

3.6 Set differentiating governance

1. For your organization, where is the strongest case that governance investment in agentic AI would create competitive advantage rather than simply protective hygiene? Name the specific market, customer segment, or regulatory environment where this would land most clearly.

Add four to six sentences here:

The strongest competitive advantage from governance investment in agentic AI lies in Huawei's enterprise and government customer segment, where data sovereignty, auditability, and compliance are non-negotiable deal requirements, not nice-to-haves. Unlike competitors competing on token price, Huawei can differentiate on trust and verifiability: a fully governed AI Project Brain with hardcoded constraints, non-human identity management, and AgentOps dashboards becomes a sales asset that proves AI can safely touch core business workflows. Governance also aligns with Huawei's stated strategy of turning project management into a "root capability". An internal strength that can be productized and offered to partners who face the same compliance-heavy, complex delivery challenges. This shifts governance from a cost of compliance to a market-differentiating capability that opens doors in regulated industries.

2. What is the credible governance posture your organization could announce externally in the next 12 months? Describe the truth behind that announcement, not the aspirational marketing.

Add three to five sentences here:

In the next 12 months, Huawei can credibly announce that its ISO/IEC 42001-certified AI governance framework has been operationalized into the AgentArts platform as default engineering guardrails, so every deployed agent inherits hardcoded constraints, identity management, and audit trails without per-project customization. It has already demonstrated 7x24 automated agent operations in regulated banking environments, and is actively participating in national agentic AI safety standards. The honest message is that governance is no longer a policy document but a platform capability. AI Project Brain may can be one of the first internal workflows to validate this in project management.

3.7 Running recommendation: The SCALE

This paragraph integrates the work of this section into a single coherent statement: the autonomy posture, the safeguard architecture, the layered lines of defense at portfolio scale, the identity discipline, and the governance-as-differentiator framing. It closes with an explicit ask: telling your board not only what you intend to do, but how you intend to do it responsibly, and what you need from them to do it well.

Optional scaffold:

[Your organization] will scale this responsibly by [autonomy posture and safeguard architecture], governed through [the layered lines of defense and identity disciplines], and treating governance as [the differentiator framing specific to your market or regulatory context]. To do this we are asking the board for [the specific investment, mandate, sponsorship, governance authority], in return for [the specific commitments your team will make].

1. Write your SCALE paragraph.

Add five to eight sentences here:

*Huawei will scale the AI Project Brain responsibly through a **disciplined, layered autonomy approach**: agents start at Delegated within hardcoded constraints, escalate to humans at defined thresholds, and migrate toward greater autonomy only when adoption rate, accuracy, and PM confidence metrics are met. The safeguard architecture rests on **four defence in-depth layers** with hardcoded constraints, output verification, reversibility via staging plus human gates, and escalation with full context, while portfolio-scale governance is enforced through a central AgentOps team, non-human identity management, and an AI Bill of Materials that tracks every model and data dependency. We position governance as a **competitive differentiator** in the enterprise and government segment, where auditability and compliance are deal-winning requirements.*

*We are asking the board for **three commitments**: formal approval to establish the AI Project Operations Specialist role via internal promotion, endorsement to*

*prioritise agent identity management on the enterprise IAM roadmap, and budget for the AgentOps dashboard and infrastructure. In return, we commit to delivering **a validated governance framework** reusable across 20 agent deployments within 12 months, reducing PM manual reporting time by 3–4 hours per week, and demonstrating an end-to-end auditable agent workflow in a regulated enterprise context, proving that governance is not a compliance burden, but a foundation for customer trust.*

2. Now read all three running recommendation paragraphs together. Does the argument hold? Is the chain of reasoning visible? Where does it strain? Make any final revisions to the earlier paragraphs you now think they need.

Add a short revision note, plus any edits you make above, here:

Add the following sentence into the first running recommendation paragraph, specifically after “This measured approach fits Huawei's current position...” to make 3 paragraphs more consistent:

“In Huawei's enterprise and government customer segment, where auditability and compliance are not optional safeguards but deal-winning requirements, this governance capability is not a compliance cost—it is a market-differentiating asset.”

Board Recommendation

This assembles your work throughout this program into an artefact you can take to your board. Though it may seem counterintuitive, build the recommendation in the order presented:

A. Body

Compile your three running recommendation paragraphs, making whatever adjustments are needed for a coherent argument.

Add your running recommendations here:

Huawei should respond to agentic AI with deliberate, strategic urgency, anchoring its initial bet on the AI Project Brain, an end-to-end autonomous workflow that spans contract management, budget control, and governance reporting. This measured approach fits Huawei's current position: in Huawei's enterprise and government customer segment, where auditability and compliance are not optional safeguards but deal-winning requirements, this governance capability is not a compliance cost, it is a market-differentiating asset; we are strong in classical ML and generative AI, but still transitioning from AI-first to AI-native, with promising pockets of native deployment (e.g., AUTINOps), yet lacking the unified data infrastructure and cross-functional governance needed for safe, large-scale agentic autonomy. The three adoption realities make this clear: agentic AI requires breaking data silos, establishing new accountability for machine errors, and fundamentally redefining human roles. By starting with a contained, high-visibility workflow that touches commercial, financial, and delivery functions without yet impacting customer-facing or network-critical systems, Huawei can safely build the organizational muscle for data sharing, governance, and role redesign, generating the evidence and trust needed to scale agentic AI confidently, without overcommitting before the cultural and infrastructural foundations are ready.

Huawei will build this capability by buying and deploying a private, on-premises LLM, stacking open-source orchestration frameworks with a self-built governance and safety wrapper, and layering a custom business abstraction layer on top of its internal systems (ISDP, MetaERP, iSales, WeLink, etc.) to shield the agent from legacy complexity. The agent will use hybrid reasoning, single-step for routine actions, chain-of-thought for complex deviations, and sub-agent decomposition for cross-functional workflows with perception spanning structured APIs and unstructured meeting/email logs, memory maintained across sessions and deployments, and safety enforced via hard financial caps, mandatory escalation routing, and immutable audit trails. Human-in-the-loop thresholds are defined by a threshold map where the agent acts autonomously for low-risk tasks, escalates to PMs for medium-risk decisions, and requires human origination for high-risk irreversible actions. The agent will be managed day-to-day by a Project Manager during the pilot, transitioning post-pilot to an AI Project Operations Specialist, an internal promotion path from experienced PMs, who will maintain thresholds, curate edge cases, and train new PMs. It will be held to standard through an AgentOps discipline with real-time metrics (recommendation adoption rate, escalation-to-action latency), monthly human audits of root-cause label accuracy, and a governance cadence where the Specialist reviews daily, the PMO Lead approves bi-weekly, and IT Security holds veto power for compliance violations.

Huawei will scale the AI Project Brain responsibly through a disciplined, layered autonomy approach: agents start at Delegated within hardcoded constraints, escalate to humans at defined thresholds, and migrate toward greater autonomy only when adoption rate, accuracy, and PM confidence metrics are met. The safeguard architecture rests on four defence in-depth layers with hardcoded constraints, output verification, reversibility via staging plus human gates, and escalation with full context, while portfolio-scale governance is enforced through a central AgentOps team, non-human identity management, and an AI Bill of Materials that tracks every model and data dependency. We position governance as a competitive differentiator in the enterprise and government segment, where auditability and compliance are deal-winning requirements. We are asking the board for three commitments: formal approval to establish the AI Project Operations Specialist role via internal promotion, endorsement to prioritise agent identity management on the enterprise IAM roadmap, and budget for the AgentOps dashboard and infrastructure. In return, we commit to delivering a validated governance framework reusable across 20 agent deployments within 12 months, reducing PM manual reporting time by 3-4 hours per week, and demonstrating an end-to-end auditable agent workflow in a regulated enterprise context, proving that governance is not a compliance burden, but a foundation for customer trust.

B1. Value creation across five dimensions

Complete each row of Table 5 with the specific (and, where possible, measurable) value your candidate workflow's agentic version would create. Be honest where the value is small, ambiguous, or speculative – your board will respect candor and suspect overclaim.

“When an agent resolves a customer issue end-to-end, the value is not only in the minutes saved, it is in the strength or weakness of the lifetime relationship that follows.”

Adapted from the Module 1 Boulton podcast

Table 5: Five dimensions.

Dimension	What this dimension is asking	Near-term value (12 months)	Long-term value (3-5 year)
Efficiency	Reduction in time, cost, or cognitive load required to complete the work and where in the workflow that	Each PM saves 3-4 hours/week on manual reporting; effort shifts from data compilation to client-facing and	Efficiency scales across the PMO; same team manages more projects without proportional

	reduction is realized	strategic work.	headcount growth. Cognitive load shifts from routine tracking to complex governance.
Quality	Improvement in accuracy, consistency, traceability, or completeness, including the kinds of error the agent makes less often than humans, and the kinds it may make more often	Consistent deviation detection (unified ≥ 5 -day threshold); full audit trail replaces scattered emails/Excel logs. Agent misses fewer deviations but may misclassify root causes due to terminology drift.	Consistent deviation detection (unified ≥ 5 -day threshold); full audit trail replaces scattered emails/Excel logs. Agent misses fewer deviations but may misclassify root causes due to terminology drift.
Resilience	Capacity to handle volume spikes, staff turnover, knowledge loss, or disruption, and what becomes more recoverable when the agent is in place	Knowledge retention: agent memory captures PM expertise (thresholds, handling logic), reducing knowledge loss during turnover. Staging & human gate prevents erroneous reports from reaching clients.	Project management capability becomes institutionalized, not dependent on a few senior PMs. New PMs onboard faster with the agent as a knowledge partner. System outages are recoverable without rework.
Speed	Reduction in end-to-end cycle time and the second-order effects of that reduction on customer	Deviation detection drops from 3–5 days to <24 hours; PMs act on deviations the next day rather than	Speed moves from daily to real-time (event-driven triggers). Faster client response becomes a

	experience, market response, or downstream processes	weekly.	differentiator in competitive bids.
Innovation	What new work becomes possible because the agent has freed capacity, opened access, or restructured the unit of effort – often the largest long-term value, and the easiest to understate	PM role begins shifting from "reporter" to "supervisor"; accumulated deviation data starts surfacing vendor level insights.	PMs evolve into project strategists, focused on client relationships, cross-project optimisation, and governance design. Agent evolves from reporter to project health advisor, proactively suggesting resource adjustments and risk mitigation. Governance framework itself may become a productised template for partners.

B2. ROI narrative: Near term and long term

The five-dimension table is the structure. The ROI narrative is the prose. Boards read tables for orientation, but they decide on the strength of the narrative that surrounds them.

1. Near-term ROI narrative (next 12 months). What will the agent demonstrably deliver in its first year, in measurable terms with the caveats your board will expect? Reference the value table where useful, but write in connected prose.

Add four to six sentences here:

In its first 12 months, the AI Project Brain will deliver three measurable outcomes. First, it will save each participating PM an estimated 3 to 4 hours per week by

automating manual reporting, and that time will be reallocated to client-facing work and strategic decision-making.

Second, it will compress deviation detection and reporting from the current 3-to-5-day cycle to under 24 hours, enabling PMs to act on project variances the next day rather than waiting for weekly or bi-weekly status updates.

Third, it will replace scattered emails and Excel logs with a full immutable audit trail for every deviation detected, every root-cause label assigned, and every PM decision. Pilot data already shows 92% alignment with human judgments, with the remaining 8% consisting largely of ambiguity cases where both outcomes are defensible.

The caveats are that root-cause label accuracy will require monthly human audits to correct terminology drift, that the agent will remain read-only by default and execute no irreversible action without a human gate, and that the productivity savings will depend on PM adoption and trust, which we will track through the AgentOps dashboard and address with targeted training.

At the end of 12 months, we will have a validated, auditable agentic workflow that reduces PM overhead, accelerates deviation response, and establishes a reusable governance model for the organisation, proving that this is not a pilot that stays a pilot but a foundation we can responsibly scale.

2. Long-term ROI narrative (three to five years). What does sustained investment in this capability, and the portfolio of agents that follows from it, make possible at the operating-model level? Be concrete: name the new units of work, the new market positions, and the new organizational capabilities.

Add four to six sentences here:

Over three to five years, sustained investment in the AI Project Brain will fundamentally reshape Huawei's project delivery operating model. PMs will evolve from task tracking to exception-based portfolio management, and new roles will emerge including AI Project Operations Specialists, AgentOps Analysts, and Project Data Stewards.

At the market level, our auditable governance framework will become a qualifying differentiator in regulated industries, opening doors with government and financial clients who require verifiable AI safety. The agent portfolio will also enable a new productised service offering, embedding governed project delivery into Huawei's broader enterprise portfolio. By year five, we will have institutionalised non-human identity management, agent-ready data pipelines, and a reusable threshold-and-escalation library, making governance a first-class engineering discipline and a durable source of competitive advantage.

C. The closing paragraph: The ask and the commitment

Every board recommendation needs a closing paragraph that brings the argument home. Remember that vague asks produce vague responses; specific asks produce decisions.

1. Write the closing paragraph of your recommendation. Name the ask (mandate, investment, sponsorship, authority); name the commitment (cadence, governance discipline, transparency, exit conditions). Close with a single sentence that names what the organization gains by acting on this now rather than later.

Add your closing paragraph in three to five sentences here:

We ask the board to approve three commitments: formal mandate to establish the AI Project Operations Specialist role via internal promotion, endorsement to prioritise non-human identity management on the enterprise IAM roadmap, and budget for the AgentOps dashboard and audit infrastructure. In return, we commit to quarterly reviews with transparent reporting on adoption, accuracy, and PM time saved, and a clear exit condition: if adoption falls below 50% or accuracy below 70% for two consecutive quarters, we will recommend pausing further deployment.

What we gain by acting now is the ability to build governance muscle on a manageable scale, rather than playing catch-up when competitive pressure forces us to move faster than we are prepared for.

D. Executive summary

Written last but placed first. The Executive Summary is the single page your board reads before they read anything else, and decides whether they read the rest. It is the compressed statement of what you are recommending, why, and on what terms.

“Boulton Industries should deploy agentic AI, deliberately and at measured pace, in a manner that treats governance, safeguard architecture, and organizational readiness as the central design variables rather than downstream constraints.”

The fictional CEO, Boulton Industries

1. Write the single opening sentence of your Executive Summary. The sentence should state, in one breath, the recommendation itself – the strategic posture, the scope, and the qualifier that signals seriousness.

Add one sentence here:

Huawei should deploy the AI Project Brain as a deliberate, governed pilot, starting with a read-only, single-project deviation detection workflow that treats governance architecture, organisational readiness, and non-human identity

management as the core design variables, not as constraints to be addressed after deployment.

2. Compress your WHAT paragraph into three to four sentences for the Summary. Strip detail; keep argument.

Add three to four sentences here:

Huawei should respond to agentic AI with deliberate, strategic urgency, anchoring its initial bet on the AI Project Brain, a governed workflow spanning contract management, budget control, and governance reporting. In Huawei's enterprise and government customer segment, where auditability and compliance are deal-winning requirements, governance is not a compliance cost but a market-differentiating asset. We are strong in classical ML and generative AI, but still transitioning from AI-first to AI-native, lacking the unified data infrastructure and cross-functional governance needed for safe, large-scale agentic autonomy. By starting with a contained, high-visibility workflow that touches commercial, financial, and delivery functions without yet impacting customer-facing or network-critical systems, Huawei can safely build the organizational muscle for data sharing, governance, and role redesign, generating the evidence and trust needed to scale agentic AI confidently, without overcommitting before the cultural and infrastructural foundations are ready.

3. Compress your HOW paragraph into three to four sentences. Name the design decisions in compressed form; do not list them.

Add three to four sentences here:

We will deploy a private, on-premises LLM wrapped in a self-built governance and safety layer, shielded from legacy systems by a custom business abstraction layer. The agent will use hybrid reasoning, single-step for routine tasks, chain-of-thought for complex deviations, and sub-agent decomposition for cross-functional workflows, with safety enforced via hard financial caps, mandatory escalation, and immutable audit trails. Autonomy is calibrated by a threshold map (autonomous for low-risk, escalated for medium-risk, human-originated for high-risk), managed daily by a PM during pilot and post-pilot by an AI Project Operations Specialist, with governance held to standard through AgentOps metrics, monthly audits, and a tri-party review cadence (Specialist daily, PMO bi-weekly, IT Security veto). The MVP intentionally excludes cross-system integration and automated execution, focusing strictly on read-only, single-project deviation detection.

4. Compress your SCALE paragraph into three to four sentences. Include the explicit ask, in one phrase.

Add three to four sentences here:

Huawei will scale the AI Project Brain through a disciplined, layered autonomy approach where agents start at Delegated within hardcoded constraints, escalate to humans at defined thresholds, and migrate toward greater autonomy only when adoption, accuracy, and PM confidence metrics are met. The safeguard architecture rests on four defence-in-depth layers, hardcoded constraints, output verification, reversibility via staging plus human gates, and escalation with full context, while portfolio-scale governance is enforced through a central AgentOps team, non-human identity management, and an AI Bill of Materials. We position governance as a competitive differentiator in the enterprise and government segment, where auditability and compliance are deal-winning requirements. We ask the board for three commitments, approval for the AI Project Operations Specialist role, prioritisation of agent identity on the IAM roadmap, and budget for the AgentOps dashboard, in return we commit to delivering a validated governance framework reusable across 20 deployments within 12 months, reducing PM manual reporting time by 3-4 hours per week, and proving that governance is not a compliance burden but a foundation for customer trust.

5. Write one closing sentence. Express what the organization gains by acting now in the most direct language you can find.

Add one sentence here:

By acting now, Huawei gains the ability to build governance muscle on a manageable scale, turning a potential compliance burden into a competitive advantage that opens doors in regulated markets before the window close.

You have now written your Executive Summary. By convention, it sits at the front of the final document, between the title page and the body. When you assemble the final version of the recommendation outside this playbook, place this Summary first, followed by the body, then the value creation table and ROI narrative, then the closing paragraph.

Module 3 rubric

	Criteria not met	Criteria met	Good	Exceptional
Sections 3.1–3.3 <i>Safeguards and autonomy</i>	No submission. OR The response fails to apply the autonomy dial and safeguard layers.	The response attempts to apply the autonomy dial and safeguard layers.	The response provides sound placement on the autonomy dial, and concrete applications of the safeguard layers.	The response provides thorough, nuanced insight into calibrating autonomy and layered safeguards. The response demonstrates a high level of expertise in the module's content.
Sections 3.4–3.6 <i>Portfolio governance</i>	No submission. OR The response does not demonstrate insight into the lines of defense or identity management.	The response attempts to provide insight into the lines of defense and identity management.	The response provides sound insight into the lines of defense and identity management.	The response provides thorough insight into the lines of defense and identity management. The response demonstrates a high level of expertise in the module's content.
Section 3.7 <i>Running recommendation: SCALE</i>	No submission. OR The response does not demonstrate insight into the running recommendation.	The response attempts to provide insight into the running recommendation.	The response provides sound insight into the running recommendation.	The response provides thorough insight into the running recommendation. The response demonstrates a high level of expertise in the module's content.
Board recommendation	No submission.	The response attempts to provide insight into the	The response provides sound insight into the	The response provides thorough insight into the

	Criteria not met	Criteria met	Good	Exceptional
<i>Value creation and ROI</i>	OR The response does not demonstrate insight into the value creation and ROI.	value creation and ROI.	value creation and ROI.	value creation and ROI. The response demonstrates a high level of expertise in the module's content.
Board recommendation <i>Closing commitment</i>	No submission. OR The closing lacks a specific ask.	The closing is vague and does not name specific board asks.	The closing provides a specific ask.	The closing provides a highly specific ask to drive the board towards immediate action. The response demonstrates a high level of expertise in the module's content.
Board recommendation <i>Executive Summary</i>	No submission. OR The response does not provide a concise and coherent Executive Summary.	The response attempts to provide a concise and coherent Executive Summary.	The response provides a concise and coherent Executive Summary.	The response provides a motivating, concise and coherent Executive Summary and demonstrates a high level of expertise in the module's content.
<i>Clarity and logic of the submission's structure</i>	No submission. OR The response is unstructured, unclear, and lacks succinctness.	The response is logically structured, clear enough to comprehend, and mostly succinct.	The response is well structured in terms of logic, clarity, and succinctness.	The response is exceptionally well structured in terms of logic, clarity, and succinctness.